

## Assignment 1 Part 1 – Regression

Q5. What conclusion if any can be drawn from the weight values? How does gender and BMI affect blood sugar levels? What are the estimated blood sugar levels for the below examples?

| AGE | SEX | BMI | BP  | S1  | S2    | S3 | S4  | S5     | S6  | y_pred   |
|-----|-----|-----|-----|-----|-------|----|-----|--------|-----|----------|
| 25  | F   | 18  | 79  | 130 | 64.8  | 61 | 2   | 4.1897 | 68  | 43.5294  |
| 50  | M   | 28  | 103 | 229 | 162.2 | 60 | 4.5 | 6.107  | 124 | 232.2309 |

Weights –

tensor ([[ 1.9400, -11.4488, 26.3047, 16.6306, -9.8810, -2.3179, -7.6995, 8.2121, 21.9769, 2.6065, 153.7365]])

**Ans**

Looking into the weights vector, weight associated with BMI is positive i.e., BMI seems to be positively co-related with blood sugar levels and as for the gender is concerned Females tend to have lower blood sugar than the males. Another noteworthy observation is that weights associated to s1-s6 are higher in comparison to the rest can be seen as major contributing factor for regression output.

Code:

```
1 x1 = torch.tensor([25, 2, 18, 79, 130, 64.8, 61, 2, 4.1897, 68]).reshape(1, -1)
2 x2 = torch.tensor([50, 1, 28, 103, 229, 162.2, 60, 4.5, 6.107, 124]).reshape(1, -1)
3
4 x1 = norm_set(x1, mu, sigma)
5 x2 = norm_set(x2, mu, sigma)
6 x1 = torch.cat([x1, torch.ones(x1.shape[0], 1)], dim=1)
7 x2 = torch.cat([x2, torch.ones(x2.shape[0], 1)], dim=1)
8
9 print(f"Blood sugar level Example 1: {model(x1)[0][0]}")
10 print(f"Blood sugar level Example 2: {model(x2)[0][0]}")
```

✓ 0.0s

Blood sugar level Example 1: 43.52943420410156  
Blood sugar level Example 2: 232.23094177246094

Q6. Try the code with several learning rates that differ by orders of magnitude and record the error of the training and test sets. What do you observe on the training error? What about the error on the test set?

Ans

Ran experiments with learning rates below –

[0.0125, 0.025, 0.05, 0.125, 0.25, 0.3, 0.45, 0.5, 0.1, 0.01, 0.001, 0.0001, 0.00001, 1, 10, 100] and saw the lowest cost of 2884.9223 for training set at learning rate of 0.05. We also observe NaNs for the learning rate with values of 1, 10, 100 as they overshoot the training error and fails to converge at the optimal cost and finally, we observe a cost of 2885.6189 on the test set, which is close to the training cost indicating a good fit.

```
... Training cost 3066.773681640625
      learning rate: 0.0125 test cost:3114.945556640625
      Training cost 2898.609619140625
      learning rate: 0.025 test cost:2890.0009765625
      Training cost 2894.800048828125
      learning rate: 0.05 test cost:2884.92236328125
      Training cost 2888.496826171875
      learning rate: 0.125 test cost:2885.8720703125
      Training cost 2907.636962890625
      learning rate: 0.25 test cost:2907.63525390625
      Training cost 10136.5478515625
      learning rate: 0.3 test cost:8.319514378800598e+30
      Training cost 15319.40234375
      learning rate: 0.45 test cost:nan
      Training cost 19150.126953125
      learning rate: 0.5 test cost:nan
      Training cost 2890.406494140625
      learning rate: 0.1 test cost:2885.618896484375
      Training cost 3356.777587890625
      learning rate: 0.01 test cost:3431.069580078125
      Training cost 20040.583984375
      learning rate: 0.001 test cost:18534.30859375
      Training cost 28468.08203125
      learning rate: 0.0001 test cost:25530.2265625
      Training cost 29583.326171875
      ...
      Training cost 29711.322265625
      learning rate: 10 test cost:nan
      Training cost 29711.322265625
      learning rate: 100 test cost:nan
```

Code:

```
1  ### your code here
2  learning_rates = [0.0125, 0.025, 0.05, 0.125, 0.25, 0.3, 0.45, 0.5, 0.1, 0.01, 0.001, 0.0001, 0.00001, 1, 10, 100]
3
4  def training_function(x_test, y_test, alpha):
5      cost_lst = list()
6      model = LinearRegression(x_train.shape[1])
7      for _ in range(100):
8          prediction = model(x_train)
9          cost = mean_squared_error(prediction, y_train)
10         cost_lst.append(cost)
11         gradient_descent_step(model, x_train, y_train, prediction, alpha)
12
13     print(f"Training cost {min(cost_lst)}")
14     test_pred = model(x_test)
15     cost = mean_squared_error(test_pred, y_test)
16     print(f"learning rate: {alpha} test cost:{cost}")
17     return cost_lst
18
19 min_costs = []
20 for alpha in learning_rates:
21     cost_lst = training_function(x_test, y_test, alpha)
22     # print(f"lr: {alpha} min cost: {min(cost_lst)}\n")
23     min_costs.append(min(cost_lst))
24
25 plt.plot(learning_rates, min_costs)
26 plt.xlabel("learning rates")
27 plt.ylabel("cost")
28 plt.show()
29
```

Q8. First, find the best value of alpha to use to optimize best. Next, experiment with different values of lambda and see how this affects the shape of the hypothesis.

Ans

We specify a range of values to search for optimal learning rate as below –

**`alphas = torch.linspace(1, 1.5, 100)`**

Keeping the regularization parameter constant (0.0001), we can train the model with the learning rate on the above values. The model converges at a cost of 0.0083 for learning rate 1.0.

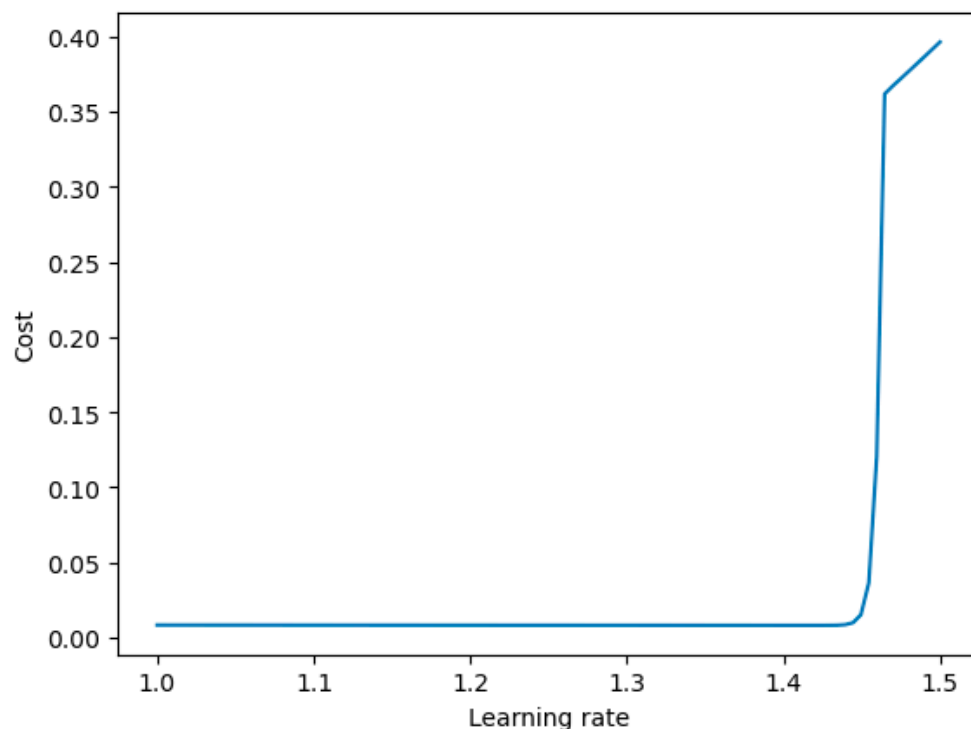
The next step would be to arrive at a value for the regularization parameter, this is achieved by setting a constant learning rate to a constant value 1.0 and running the model for regularization parameter for the below values –

**`lambdas = torch.linspace(0, 1, 10000)[1:-1]`**

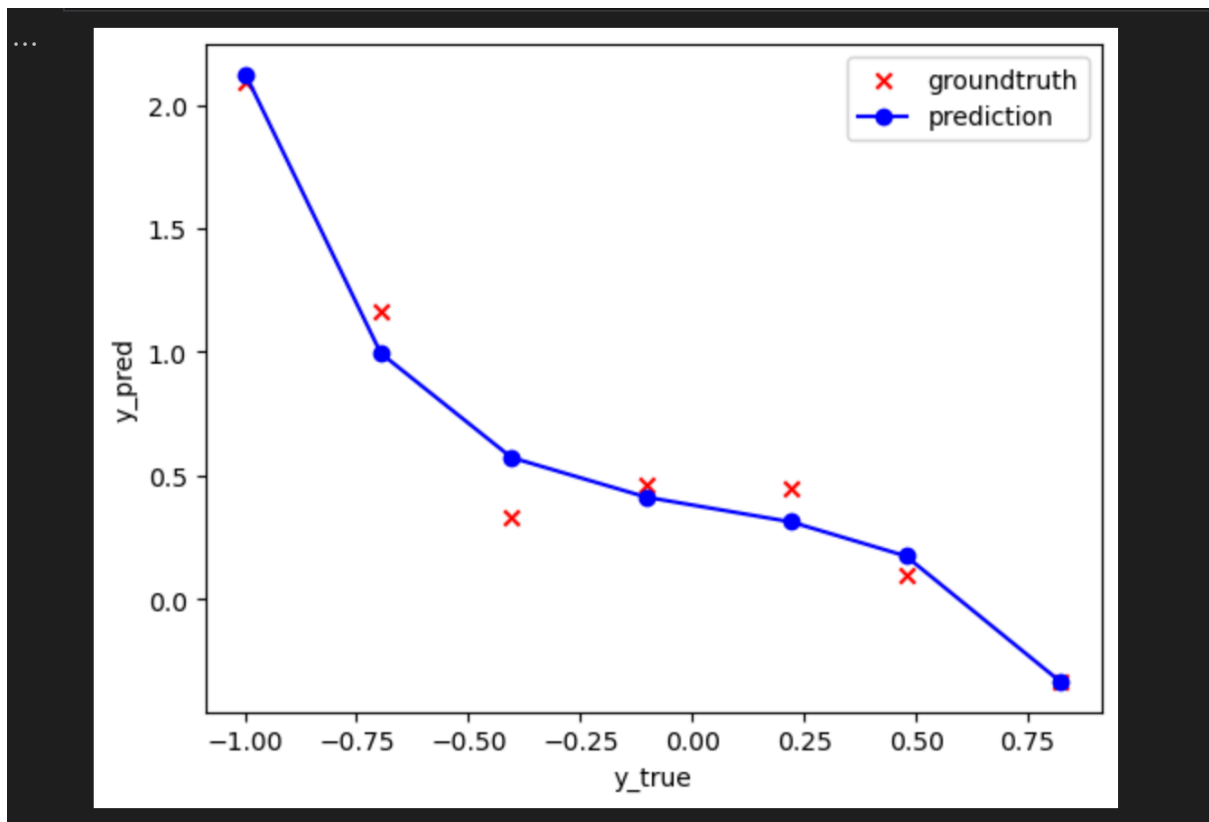
We notice that, keeping the learning rate constant 1.0, we can train the model with the regularization parameter on the above values. The model converges at a cost of 0.0083 for learning rate 0.0001.

Thus, we can conclude this model has the low cost at 0.0083 by training with the hyperparameters, learning rate set at 1.0 and the regularization param set at 0.0001.

The below graph indicates that learning rates above 1.45 tends to diverge from the hypothesis function.



The visualization below confirms that the above model creates a good fit for the data (albeit limited data) with minimal errors with the ground truth for learning rate 1.0 and regularisation param 0.0001.



For the below values of regularisation param-  
 Lambdas = [0.00001, 0.0001, 0.001, 0.01, 0.1, 0, 1, 5, 10, 100], the model underfits starting  
 from 2 as the regularization parameter now has a high value it tends to penalizes the cost  
 function much more causing it to underfit.

